

Type *Markdown* and LaTeX: α^2

Questions –

1. Find out the make and model that is most represented and the one that is least represented class? Is there a class imbalance problem, how would you handle class imbalance if it were to exist (2 marks)
2. Apply the image processing techniques and explain the benefits of those technique (2 marks)
3. Train a lightweight model, show train and validation loss curves, show steps taken to tune hyper params (1 marks)
4. Determine metric to evaluate performance of the model. Report how the model is performing on the metric (2 Marks)
5. Deploy (3 marks) a. As a rest api endpoint b. Mobil app model

```
1 # Car Images classification using CNN
```

```
1 # Import Modules
```

In [8]:

```
1 # import the libraries as shown below
2 import os
3 import shutil
4 import random
5 import numpy as np
6 import numpy as np
7 from glob import glob
8 import tensorflow as tf
9 import matplotlib.pyplot as plt
10 from tensorflow.keras.models import Model
11 from tensorflow.keras.optimizers import Adam
12 from tensorflow.keras.models import Sequential
13 from tensorflow.keras.preprocessing import image
14 from tensorflow.keras.applications.resnet50 import ResNet50
15 from tensorflow.keras.applications.resnet50 import preprocess_input
16 from tensorflow.keras.layers import Input, Lambda, Dense, Flatten, Dropout
17 from tensorflow.keras.preprocessing.image import ImageDataGenerator, load_img
18 from tensorflow.keras.callbacks import LearningRateScheduler, EarlyStopping
19 from tensorflow.keras.preprocessing.image import ImageDataGenerator, load_img, img_
```

In [4]:

```
1 import tensorflow as tf
2 print("TensorFlow version:", tf.__version__)
```

TensorFlow version: 2.4.1

In [18]:

```
1 #Set all the Constants
2 BATCH_SIZE = 32
3 IMAGE_SIZE = 224
4 CHANNELS=3
5 EPOCHS=50
```

In [6]:

```
1 #defining the path
2 train_path = 'train'
3 valid_path = 'test'
4
```

1) Find out the make and model that is most represented and the one that is least represented class?

In [8]:

```
1 data_dir = train_path
2 class_counts = {}
3
4 for class_name in os.listdir(data_dir):
5     class_dir = os.path.join(data_dir, class_name)
6     if os.path.isdir(class_dir):
7         class_count = len(os.listdir(class_dir))
8         class_counts[class_name] = class_count
9
10 # Find the most and least represented classes
11 most_represented_class = max(class_counts, key=class_counts.get)
12 least_represented_class = min(class_counts, key=class_counts.get)
13
14 print(f"Most Represented Class: {most_represented_class}, Count: {class_counts[most_represented_class]}")
15 print(f"Least Represented Class: {least_represented_class}, Count: {class_counts[least_represented_class]}")
```

Most Represented Class: Audi, Count: 814

Least Represented Class: Hyundai Creta, Count: 271

Is there a class imbalance problem, how would you handle class imbalance if it were to exist Apply the image processing techniques and explain the benefits of those technique

In [9]:

```
1 train_folder = train_path
2 # Initialize an empty dictionary to store class counts
3 class_counts = {}
4
5 # Iterate through the subfolders (each representing a class)
6 for class_folder in os.listdir(train_folder):
7     if os.path.isdir(os.path.join(train_folder, class_folder)):
8         # Count the number of files (images) in each class folder
9         class_count = len(os.listdir(os.path.join(train_folder, class_folder)))
10        # Store the class count in the dictionary with the class name as the key
11        class_counts[class_folder] = class_count
12
13 # Print the class distribution for the training set
14 for car_class, count in class_counts.items():
15     print(f"Class: {car_class}, Count: {count}")
16
```

Class: Hyundai Creta, Count: 271
Class: Mahindra Scorpio, Count: 316
Class: Swift, Count: 424
Class: Rolls Royce, Count: 311
Class: Tata Safari, Count: 441
Class: Toyota Innova, Count: 775
Class: Audi, Count: 814

In [13]:

```

1  #balance dataset
2  # Define the path to your train data directory
3  train_data_dir = 'train' # Update with your actual path
4
5  # Create a directory for balanced training data
6  balanced_train_dir = 'balanced_train_dir'
7  os.makedirs(balanced_train_dir, exist_ok=True)
8  Create an instance of the ImageDataGenerator with desired transformations
9  datagen = ImageDataGenerator(
10     rescale=1./255,           # Normalize pixel values to the range [0, 1]
11     rotation_range=20,        # Randomly rotate images by up to 20 degrees
12     width_shift_range=0.2,     # Randomly shift the width of images by up to 20%
13     height_shift_range=0.2,   # Randomly shift the height of images by up to 20%
14     shear_range=0.2,          # Randomly apply shear transformations
15     zoom_range=0.2,           # Randomly zoom in on images by up to 20%
16     horizontal_flip=True,     # Randomly flip images horizontally
17     fill_mode='nearest'       # Fill in empty areas created by transformations
18 )
19
20 # Define the desired count for 'Audi'
21 desired_count = 814
22
23 # List all class directories and their counts
24 class_directories = os.listdir(train_data_dir)
25 class_counts = {}
26
27 # Iterate through each class directory
28 for class_name in class_directories:
29     class_dir = os.path.join(train_data_dir, class_name)
30     class_files = os.listdir(class_dir)
31
32     # Create a directory for the class in the balanced dataset
33     balanced_class_dir = os.path.join(balanced_train_dir, class_name)
34     os.makedirs(balanced_class_dir, exist_ok=True)
35
36     # Calculate the number of samples needed for this class to match the desired count
37     num_original_samples = len(class_files)
38     num_samples_to_copy = min(desired_count, num_original_samples)
39
40     # Calculate the number of samples needed to augment
41     num_samples_to_augment = desired_count - num_samples_to_copy
42
43     # Copy the original samples to the balanced dataset
44     files_to_copy = random.sample(class_files, num_samples_to_copy)
45     for file_name in files_to_copy:
46         src_path = os.path.join(class_dir, file_name)
47         dst_path = os.path.join(balanced_class_dir, file_name)
48         shutil.copy(src_path, dst_path)
49
50     # Update the class counts
51     class_counts[class_name] = num_samples_to_copy
52
53     # Apply data augmentation if needed
54     if num_samples_to_augment > 0:
55         # Iterate through the original samples and apply augmentation
56         for i in range(num_samples_to_augment):
57             # Randomly select an image from the original samples
58             original_image_name = random.choice(class_files)
59             original_image_path = os.path.join(class_dir, original_image_name)

```

```
60
61     # Load the image using PIL (Pillow)
62     img = load_img(original_image_path)
63
64     # Convert the image to a NumPy array
65     img_array = img_to_array(img)
66
67     # Reshape the image to (1, height, width, channels)
68     img_array = np.expand_dims(img_array, axis=0)
69
70     # Generate augmented images using the data generator
71     augmented_images = datagen.flow(img_array, batch_size=1)
72
73     # Get the first augmented image from the generator
74     augmented_image_array = next(augmented_images)[0].astype(np.uint8)
75     augmented_image = array_to_img(augmented_image_array)
76
77     # Save the augmented image with a new name
78     augmented_image_name = f"augmented_{i}_{original_image_name}"
79     augmented_image_path = os.path.join(balanced_class_dir, augmented_image_name)
80     augmented_image.save(augmented_image_path)
81
82
83
```

Rescaling (Normalization):

Benefit: Normalizing pixel values to the range [0, 1] helps in convergence during training by ensuring that the model deals with small and consistent input values. Rotation, Width Shift, Height Shift, Shear, and Zoom:

Benefit: These transformations increase the diversity of your training data. They help the model become robust to variations in input images, such as changes in orientation, position, and scale. Horizontal Flip:

Benefit: Flipping images horizontally introduces mirror-image versions of existing data, further enhancing dataset diversity and training stability. Fill Mode (Nearest):

Benefit: This mode fills in areas of the image that may become empty due to transformations like rotation or shearing. It ensures that no information is lost during augmentation.

In [14]:

```
1 #after blanching dataset
```

In [10]:

```
1 train_folder = 'balanced_train_dir'
2 # Initialize an empty dictionary to store class counts
3 class_counts = {}
4
5 # Iterate through the subfolders (each representing a class)
6 for class_folder in os.listdir(train_folder):
7     if os.path.isdir(os.path.join(train_folder, class_folder)):
8         # Count the number of files (images) in each class folder
9         class_count = len(os.listdir(os.path.join(train_folder, class_folder)))
10        # Store the class count in the dictionary with the class name as the key
11        class_counts[class_folder] = class_count
12
13 # Print the class distribution for the training set
14 for car_class, count in class_counts.items():
15     print(f"Class: {car_class}, Count: {count}")
```

Class: Hyundai Creta, Count: 814
Class: Mahindra Scorpio, Count: 814
Class: Swift, Count: 814
Class: Rolls Royce, Count: 814
Class: Tata Safari, Count: 814
Class: Toyota Innova, Count: 814
Class: Audi, Count: 814

Train a lightweight model, show train and validation loss curves, show steps taken to tune hyper params

In [19]:

```
1 dataset = tf.keras.preprocessing.image_dataset_from_directory(
2     "balanced_train_dir",
3     seed=123,
4     shuffle=True,
5     image_size=(IMAGE_SIZE, IMAGE_SIZE),
6     batch_size=BATCH_SIZE
7 )
```

Found 5698 files belonging to 7 classes.

In [20]:

```
1 class_names = dataset.class_names
2 print(class_names)
3 len(class_names)
```

['Audi', 'Hyundai Creta', 'Mahindra Scorpio', 'Rolls Royce', 'Swift', 'Tata Safari', 'Toyota Innova']

Out[20]:

7

In [17]:

```
1 len(dataset)
```

Out[17]:

179

In [79]:

```
1 class_labels = training_set.class_indices
2 print(class_labels)
```

```
{'Audi': 0, 'Hyundai Creta': 1, 'Mahindra Scorpio': 2, 'Rolls Royce': 3,
'Swift': 4, 'Tata Safari': 5, 'Toyota Innova': 6}
```

In [18]:

```
1 class_labels = np.array(dataset.class_names)
2 print(class_labels)
```

```
['Audi' 'Hyundai Creta' 'Mahindra Scorpio' 'Rolls Royce' 'Swift'
'Tata Safari' 'Toyota Innova']
```

In [80]:

```
1 class_labels_list = list(class_labels.keys())
2 print(class_labels_list)
3
```

```
['Audi', 'Hyundai Creta', 'Mahindra Scorpio', 'Rolls Royce', 'Swift', 'Ta
ta Safari', 'Toyota Innova']
```

```
1 # Visualize Dataset
```


In [19]:

```

1 #plotting somesamples
2 plt.figure(figsize=(18, 18))
3 for image_batch, labels_batch in dataset.take(1):
4     for i in range(12):
5         ax = plt.subplot(3, 4, i + 1)
6         plt.imshow(image_batch[i].numpy().astype("uint8"))
7         plt.title(class_names[labels_batch[i]])
8         plt.axis("off")

```

2023-09-10 02:01:51.512802: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:116] None of the MLIR optimization passes are enabled (registered 2)

2023-09-10 02:01:51.529881: I tensorflow/core/platform/profile_utils/cpu_utils.cc:112] CPU Frequency: 2599990000 Hz

Swift



Mahindra Scorpio



Mahindra Scorpio



Rolls Royce



Hyundai Creta



Mahindra Scorpio



Audi



Swift



Mahindra Scorpio



Swift



Toyota Innova



Mahindra Scorpio



1 **# Create Model**

In [20]:

```
1 #importing model
2 # Here we will be using imagenet weights
3 IMAGE_SIZE = [224, 224]
4
5 resnet = ResNet50(input_shape=IMAGE_SIZE + [3], weights='imagenet', include_top=False)
6
```

In [21]:

```
1 # don't train existing weights
2 for layer in resnet.layers:
3     layer.trainable = False
```

In [22]:

```
1 # useful for getting number of output classes
2 folders = glob("balanced_train_dir/*")
```

In [23]:

```
1 # our layers
2 x = Flatten()(resnet.output)
```

In [24]:

```
1 prediction = Dense(len(folders), activation='softmax')(x)
2
3 # create a model object
4 model = Model(inputs=resnet.input, outputs=prediction)
```

In [25]:

```

1
2 # view the structure of the model
3 model.summary()
4

```

Model: "model"

Layer (type) connected to	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 224, 224, 3)]	0	
conv1_pad (ZeroPadding2D) input_1[0][0]	(None, 230, 230, 3)	0	input_1
conv1_conv (Conv2D) conv1_pad[0][0]	(None, 112, 112, 64)	9472	conv1_pad
conv1_bn (BatchNormalization) conv1_conv[0][0]	(None, 112, 112, 64)	256	conv1_conv

In [26]:

```

1 # Compiling the model
2 model.compile(
3     loss='categorical_crossentropy',
4     optimizer='adam',
5     metrics=['accuracy']
6 )
7

```

In [10]:

```

1 # Use the Image Data Generator to import the images from the dataset
2 from tensorflow.keras.preprocessing.image import ImageDataGenerator
3
4 train_datagen = ImageDataGenerator(rescale = 1./255,)
5
6 test_datagen = ImageDataGenerator(rescale = 1./255)

```

In [11]:

```

1 # Make sure you provide the same target size as initialied for the image size
2 training_set = train_datagen.flow_from_directory('balanced_train_dir',
3                                                  target_size = (224, 224),
4                                                  batch_size = 32,
5                                                  class_mode = 'categorical')

```

Found 5698 images belonging to 7 classes.

In [12]:

```
1 test_set = test_datagen.flow_from_directory('test',
2                                           target_size = (224, 224),
3                                           batch_size = 32,
4                                           class_mode = 'categorical')
```

Found 813 images belonging to 7 classes.

In [40]:

```
1
2 # Define a Learning rate schedule
3 def lr_schedule(epoch):
4     if epoch < 10:
5         return 0.001
6     elif epoch < 20:
7         return 0.0001
8     else:
9         return 0.00001
10
11 # Define early stopping
12 early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
13
14
15
```

In [42]:


```
1 # fit the model
2 # Run the cell. It will take some time to execute
3 # Train the model with hyperparameter tuning
4 r = model.fit_generator(
5     training_set,
6     validation_data=test_set,
7     epochs=50,
8     steps_per_epoch=len(training_set),
9     validation_steps=len(test_set),
10     callbacks=[LearningRateScheduler(lr_schedule), early_stopping]
11 )
```

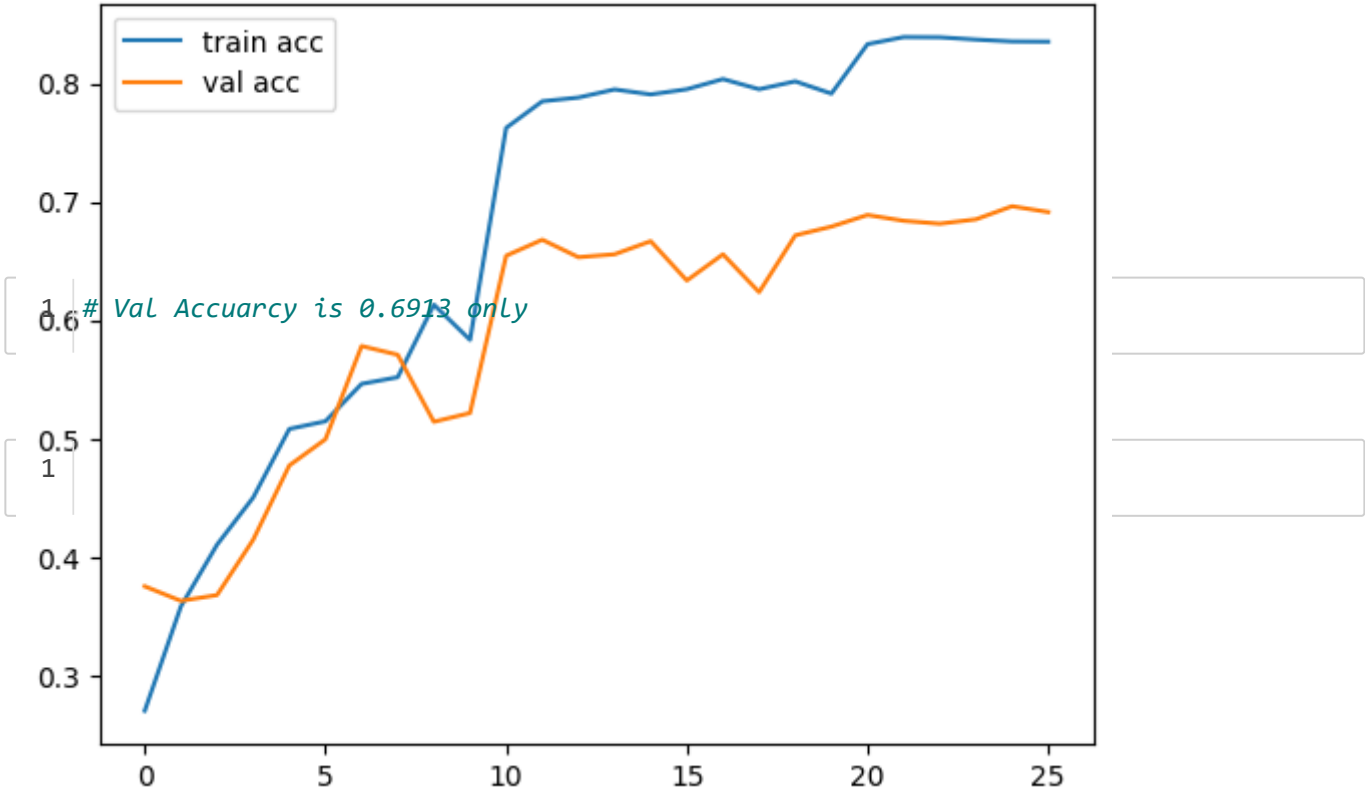
```
Epoch 1/50
179/179 [=====] - 210s 1s/step - loss: 2.6638 -
accuracy: 0.2701 - val_loss: 2.0943 - val_accuracy: 0.3752
Epoch 2/50
179/179 [=====] - 226s 1s/step - loss: 2.0530 -
accuracy: 0.3582 - val_loss: 1.9982 - val_accuracy: 0.3629
Epoch 3/50
179/179 [=====] - 205s 1s/step - loss: 1.8957 -
accuracy: 0.4103 - val_loss: 2.2804 - val_accuracy: 0.3678
Epoch 4/50
179/179 [=====] - 203s 1s/step - loss: 1.7882 -
accuracy: 0.4502 - val_loss: 2.0547 - val_accuracy: 0.4145
Epoch 5/50
179/179 [=====] - 207s 1s/step - loss: 1.4617 -
accuracy: 0.5081 - val_loss: 1.6201 - val_accuracy: 0.4772
Epoch 6/50
179/179 [=====] - 207s 1s/step - loss: 1.5643 -
accuracy: 0.5146 - val_loss: 1.5920 - val_accuracy: 0.4994
Epoch 7/50
179/179 [=====] - 203s 1s/step - loss: 1.3692 -
accuracy: 0.5462 - val_loss: 1.3316 - val_accuracy: 0.5781
Epoch 8/50
179/179 [=====] - 209s 1s/step - loss: 1.4686 -
accuracy: 0.5518 - val_loss: 1.3376 - val_accuracy: 0.5707
Epoch 9/50
179/179 [=====] - 227s 1s/step - loss: 1.1293 -
accuracy: 0.6130 - val_loss: 1.4761 - val_accuracy: 0.5141
Epoch 10/50
179/179 [=====] - 221s 1s/step - loss: 1.2427 -
accuracy: 0.5834 - val_loss: 1.9673 - val_accuracy: 0.5215
Epoch 11/50
179/179 [=====] - 220s 1s/step - loss: 0.7200 -
accuracy: 0.7624 - val_loss: 1.1035 - val_accuracy: 0.6544
Epoch 12/50
179/179 [=====] - 219s 1s/step - loss: 0.6712 -
accuracy: 0.7848 - val_loss: 1.0814 - val_accuracy: 0.6679
Epoch 13/50
179/179 [=====] - 218s 1s/step - loss: 0.6654 -
accuracy: 0.7880 - val_loss: 1.0721 - val_accuracy: 0.6531
Epoch 14/50
179/179 [=====] - 217s 1s/step - loss: 0.6607 -
accuracy: 0.7948 - val_loss: 1.1252 - val_accuracy: 0.6556
Epoch 15/50
179/179 [=====] - 217s 1s/step - loss: 0.6600 -
accuracy: 0.7906 - val_loss: 1.0647 - val_accuracy: 0.6667
Epoch 16/50
179/179 [=====] - 216s 1s/step - loss: 0.6472 -
accuracy: 0.7950 - val_loss: 1.1569 - val_accuracy: 0.6335
Epoch 17/50
179/179 [=====] - 216s 1s/step - loss: 0.6460 -
accuracy: 0.8036 - val_loss: 1.0726 - val_accuracy: 0.6556
Epoch 18/50
179/179 [=====] - 211s 1s/step - loss: 0.6467 -
accuracy: 0.7950 - val_loss: 1.1633 - val_accuracy: 0.6236
Epoch 19/50
179/179 [=====] - 223s 1s/step - loss: 0.6380 -
accuracy: 0.8017 - val_loss: 1.0764 - val_accuracy: 0.6716
Epoch 20/50
179/179 [=====] - 224s 1s/step - loss: 0.6555 -
accuracy: 0.7913 - val_loss: 1.0393 - val_accuracy: 0.6790
Epoch 21/50
```

```
179/179 [=====] - 220s 1s/step - loss: 0.5938 -  
accuracy: 0.8331 - val_loss: 1.0051 - val_accuracy: 0.6888  
Epoch 22/50  
179/179 [=====] - 234s 1s/step - loss: 0.5909 -  
accuracy: 0.8391 - val_loss: 1.0115 - val_accuracy: 0.6839  
Epoch 23/50  
179/179 [=====] - 222s 1s/step - loss: 0.5912 -  
accuracy: 0.8389 - val_loss: 1.0268 - val_accuracy: 0.6814  
Epoch 24/50  
179/179 [=====] - 224s 1s/step - loss: 0.5891 -  
accuracy: 0.8370 - val_loss: 1.0250 - val_accuracy: 0.6851  
Epoch 25/50  
179/179 [=====] - 215s 1s/step - loss: 0.5903 -  
accuracy: 0.8354 - val_loss: 1.0170 - val_accuracy: 0.6962  
Epoch 26/50  
179/179 [=====] - 220s 1s/step - loss: 0.5888 -  
accuracy: 0.8352 - val_loss: 1.0156 - val_accuracy: 0.6913
```

In [43]:

```
1 # plot the loss
2 plt.plot(r.history['loss'], label='train loss')
3 plt.plot(r.history['val_loss'], label='val loss')
4 plt.legend()
5 plt.show()
6 plt.savefig('LossVal_loss')
7
8 # plot the accuracy
9 plt.plot(r.history['accuracy'], label='train acc')
10 plt.plot(r.history['val_accuracy'], label='val acc')
11 plt.legend()
12 plt.show()
13 plt.savefig('AccVal_acc')
```





<Figure size 640x480 with 0 Axes>

In [45]:

```
1  # Define the image size, batch size, and other hyperparameters
2  IMAGE_SIZE = (224, 224)
3  BATCH_SIZE = 32
4  EPOCHS = 50
5
6  # Define early stopping
7  early_stopping = EarlyStopping(monitor='val_loss', patience=10, restore_best_weight
8
9  # Create a data generator with aggressive data augmentation for training
10 train_datagen = ImageDataGenerator(
11     rescale=1./255,
12     rotation_range=30,      # Increased rotation range
13     width_shift_range=0.3,  # Increased width shift range
14     height_shift_range=0.3, # Increased height shift range
15     shear_range=0.2,
16     zoom_range=0.2,
17     horizontal_flip=True,
18     fill_mode='nearest'
19 )
20
21 # Create a data generator for validation (no augmentation)
22 validation_datagen = ImageDataGenerator(rescale=1./255)
23
24 # Load and prepare the training data
25 train_generator = train_datagen.flow_from_directory(
26     'balanced_train_dir',
27     target_size=IMAGE_SIZE,
28     batch_size=BATCH_SIZE,
29     class_mode='categorical'
30 )
31
32 # Load and prepare the validation data
33 validation_generator = validation_datagen.flow_from_directory(
34     'test',
35     target_size=IMAGE_SIZE,
36     batch_size=BATCH_SIZE,
37     class_mode='categorical'
38 )
39
```

Found 5698 images belonging to 7 classes.

Found 813 images belonging to 7 classes.

In [50]:

```

1
2 # Load the pre-trained ResNet50 model
3 resnet = ResNet50(input_shape=IMAGE_SIZE + (3,), weights='imagenet', include_top=False)
4
5 # Fine-tune some of the later layers of the ResNet50 base
6 for layer in resnet.layers[:-10]: # Fine-tuning last 10 layers
7     layer.trainable = True
8
9 # Add custom classification layers with dropout
10 x = Flatten()(resnet.output)
11 x = Dropout(0.5)(x) # Added dropout layer
12 predictions = Dense(len(train_generator.class_indices), activation='softmax')(x)
13
14 # Create the model
15 model = Model(inputs=resnet.input, outputs=predictions)
16
17 # Compile the model with a lower learning rate
18 model.compile(
19     loss='categorical_crossentropy',
20     optimizer=Adam(lr=0.0001), # Adjusted learning rate
21     metrics=['accuracy']
22 )
23

```

In [51]:

```

1 # Train the model with increased batch size
2 history = model.fit(
3     train_generator,
4     epochs=EPOCHS,
5     steps_per_epoch=len(train_generator),
6     validation_data=validation_generator,
7     validation_steps=len(validation_generator),
8     callbacks=[LearningRateScheduler(lr_schedule), early_stopping]
9 )

```

Epoch 1/50

179/179 [=====] - 942s 5s/step - loss: 5.3101
 - accuracy: 0.1837 - val_loss: 206.5125 - val_accuracy: 0.0923

Epoch 2/50

179/179 [=====] - 934s 5s/step - loss: 2.2589
 - accuracy: 0.1859 - val_loss: 3.1113 - val_accuracy: 0.0923

Epoch 3/50

179/179 [=====] - 914s 5s/step - loss: 1.8233
 - accuracy: 0.2676 - val_loss: 2.1553 - val_accuracy: 0.0923

Epoch 4/50

179/179 [=====] - 918s 5s/step - loss: 1.7083
 - accuracy: 0.3314 - val_loss: 3.2394 - val_accuracy: 0.0923

Epoch 5/50

179/179 [=====] - 939s 5s/step - loss: 1.5959
 - accuracy: 0.4053 - val_loss: 3.3886 - val_accuracy: 0.0824

Epoch 6/50

179/179 [=====] - 966s 5s/step - loss: 1.4248
 - accuracy: 0.4771 - val_loss: 1.6765 - val_accuracy: 0.3444

Epoch 7/50

179/179 [=====] - 916s 5s/step - loss: 1.3065


In [52]:


```
1 # save it as a h5 file
2
3
4 from tensorflow.keras.models import load_model
5
6 model.save('model_resnet50.h5')
```

```
1 # Plotting the accuracy and loss curves
```

In [81]:

```
1 # plot the loss
2 plt.plot(history.history['loss'], label='train loss')
3 plt.plot(history.history['val_loss'], label='val loss')
4 plt.legend()
5 plt.show()
6 plt.savefig('LossVal_loss')
7
8 # plot the accuracy
9 plt.plot(history.history['accuracy'], label='train acc')
10 plt.plot(history.history['val_accuracy'], label='val acc')
11 plt.legend()
12 plt.show()
13 plt.savefig('AccVal_acc')
```





These accuracy values are an indication of how well the model is performing on the classification task. The higher the accuracy, the better the model is at correctly classifying car images. It's important to look at both training and validation accuracy to assess whether the model is overfitting (i.e., performing well on the training data but poorly on unseen data).

the validation accuracy is close to the training accuracy, suggesting that the model is generalizing well to unseen data.

To summarize, the model is performing well on the classification task, achieving an accuracy of around 92-94% on both the training and validation datasets.

In []:

```
1 #Predictations test data
```

In [54]:

```
1
2 y_pred = model.predict(test_set)
3
```

In [55]:

```
1 y_pred
```

Out[55]:

```
array([[4.9484480e-04, 1.6264601e-05, 4.9876604e-08, ..., 4.3519353e-07,
        1.8654418e-08, 9.9948823e-01],
       [9.9923325e-01, 3.1886596e-08, 1.6346821e-05, ..., 1.4597697e-07,
        1.2856311e-06, 1.6910975e-06],
       [9.0483147e-05, 1.0181155e-04, 9.9973267e-01, ..., 5.3417116e-06,
        1.3362401e-05, 4.4991310e-05],
       ...,
       [1.8275303e-09, 2.7658618e-09, 1.4352742e-13, ..., 1.0000000e+00,
        1.0588677e-10, 5.7752483e-09],
       [8.6054235e-05, 4.1282263e-05, 9.6249998e-08, ..., 9.3085415e-05,
        9.9977869e-01, 1.6830202e-07],
       [3.0426646e-03, 1.5796209e-09, 2.5382731e-06, ..., 5.1020194e-08,
        7.3144967e-08, 1.0517252e-08]], dtype=float32)
```

In [56]:

```
1 import numpy as np
2 y_pred = np.argmax(y_pred, axis=1)
```

In [57]:

```
1 y_pred
```

Out[57]:

```
array([6, 0, 2, 5, 0, 1, 6, 3, 1, 0, 5, 0, 0, 1, 4, 5, 0, 6, 0, 0, 0, 6,
       0, 5, 3, 0, 0, 5, 6, 4, 3, 2, 1, 0, 0, 4, 3, 1, 0, 4, 4, 3, 1, 0,
       6, 0, 6, 1, 2, 0, 6, 1, 0, 6, 2, 0, 4, 0, 4, 6, 5, 2, 5, 6, 5, 0,
       0, 0, 6, 4, 3, 0, 2, 5, 2, 6, 0, 0, 2, 2, 6, 5, 2, 5, 4, 5, 6, 6,
       0, 3, 6, 5, 3, 0, 4, 0, 5, 1, 2, 5, 0, 6, 5, 4, 4, 0, 6, 0, 0, 4,
       0, 2, 6, 4, 0, 3, 4, 6, 5, 6, 0, 5, 1, 6, 0, 6, 0, 5, 0, 5, 3, 4,
       3, 3, 2, 6, 6, 0, 4, 3, 6, 0, 0, 0, 4, 1, 6, 5, 6, 0, 4, 4, 3, 4,
       3, 5, 6, 3, 5, 2, 5, 5, 5, 6, 0, 6, 0, 2, 4, 6, 5, 2, 6, 3, 0, 5,
       3, 6, 0, 5, 0, 3, 6, 6, 6, 6, 0, 6, 5, 0, 4, 6, 4, 6, 3, 2, 0, 5,
       0, 6, 6, 6, 5, 4, 4, 6, 6, 6, 0, 6, 4, 0, 4, 6, 5, 5, 6, 0, 0, 4,
       0, 3, 0, 0, 3, 5, 5, 3, 5, 6, 6, 0, 6, 2, 1, 4, 0, 4, 6, 6, 6, 3,
       2, 6, 0, 6, 6, 5, 2, 0, 0, 6, 0, 6, 6, 3, 3, 0, 0, 2, 3, 0, 6, 5,
       2, 4, 3, 0, 6, 1, 0, 6, 5, 6, 0, 0, 4, 6, 2, 3, 5, 0, 6, 4, 0, 2,
       5, 4, 0, 3, 1, 0, 6, 0, 1, 0, 3, 2, 2, 1, 4, 5, 6, 5, 6, 0, 3, 6,
       5, 6, 4, 2, 5, 0, 0, 2, 4, 0, 6, 3, 3, 0, 3, 2, 4, 1, 0, 5, 0, 4,
       4, 6, 6, 3, 6, 0, 6, 1, 6, 0, 4, 0, 3, 0, 4, 6, 0, 6, 3, 6, 3, 0,
       2, 0, 6, 0, 5, 0, 6, 2, 0, 6, 0, 6, 0, 1, 2, 3, 5, 3, 1, 0, 6, 6,
       6, 2, 4, 4, 6, 5, 0, 6, 4, 2, 0, 0, 0, 1, 6, 5, 6, 5, 0, 6, 0, 6,
       2, 5, 0, 4, 5, 5, 0, 6, 6, 3, 5, 5, 6, 0, 0, 5, 3, 5, 0, 2, 2, 4,
       6, 4, 0, 0, 5, 4, 4, 6, 2, 0, 6, 5, 5, 0, 4, 2, 2, 4, 3, 6, 4, 2,
       1, 4, 0, 2, 0, 0, 6, 0, 6, 4, 0, 4, 0, 2, 0, 6, 6, 2, 4, 6, 6, 1,
       5, 1, 0, 0, 5, 1, 2, 3, 5, 6, 5, 1, 3, 4, 6, 6, 4, 5, 0, 0, 6, 2,
       0, 3, 0, 0, 1, 0, 6, 6, 0, 6, 2, 0, 3, 6, 5, 5, 3, 3, 6, 5, 5, 6,
       5, 6, 0, 0, 0, 4, 2, 6, 0, 0, 4, 4, 1, 2, 5, 4, 0, 3, 1, 0, 3, 0,
       3, 5, 1, 0, 5, 3, 5, 3, 6, 3, 1, 0, 6, 5, 0, 1, 6, 6, 6, 6, 4, 6,
       2, 5, 0, 1, 6, 0, 0, 1, 3, 0, 0, 0, 0, 1, 6, 2, 6, 6, 5, 0, 5, 4,
       0, 4, 2, 5, 5, 1, 3, 0, 6, 5, 2, 6, 6, 6, 6, 0, 6, 0, 5, 0, 6, 5,
       2, 0, 5, 6, 6, 6, 6, 1, 0, 6, 0, 6, 0, 0, 3, 4, 0, 6, 5, 6, 6, 2,
       0, 3, 6, 1, 6, 4, 5, 0, 5, 2, 0, 6, 4, 6, 0, 4, 3, 1, 0, 1, 5, 0,
       6, 2, 0, 6, 5, 4, 2, 0, 3, 1, 1, 0, 6, 5, 6, 0, 6, 5, 0, 4, 1, 2,
       0, 6, 0, 0, 5, 0, 6, 4, 4, 6, 0, 0, 4, 0, 4, 6, 0, 0, 4, 0, 0, 6,
       0, 6, 1, 3, 5, 0, 0, 6, 4, 1, 5, 6, 0, 2, 6, 0, 0, 2, 0, 4, 6, 3,
       4, 5, 0, 6, 4, 1, 6, 6, 4, 6, 3, 6, 4, 6, 2, 0, 3, 0, 1, 6, 2, 5,
       6, 1, 5, 3, 0, 0, 1, 0, 6, 1, 5, 5, 6, 4, 0, 0, 3, 6, 6, 6, 6, 6,
       0, 6, 4, 0, 5, 0, 1, 5, 0, 6, 1, 2, 4, 6, 6, 3, 2, 3, 0, 0, 6, 0,
       4, 4, 3, 4, 6, 4, 0, 0, 0, 3, 1, 3, 5, 0, 3, 4, 1, 5, 5, 0, 6, 6,
       4, 3, 4, 3, 6, 0, 0, 4, 4, 6, 0, 5, 0, 5, 5, 4, 2, 3, 4, 5, 3])
```

In []:

```
1 #testing the model
```

In [4]:

```
1 from tensorflow.keras.models import load_model
2 from tensorflow.keras.preprocessing import image
```


In [5]:

```
1 model=load_model('model_resnet50.h5')
```

In [6]:

```
1 model
```

Out[6]:

```
<keras.engine.functional.Functional at 0x1f47750bc40>
```

In [51]:

```
1 img=image.load_img('test/Audi/23.jpg',target_size=(224,224))
2
3
```

In [55]:

```
1 x=image.img_to_array(img)
2
```

In [56]:

```
1 x.shape
```

Out[56]:

```
(224, 224, 3)
```

In [57]:

```
1 x=x/255
```

In [47]:

```
1 x=np.expand_dims(x,axis=0)
2 #img_data=preprocess_input(x)
3 img_data.shape
```

Out[47]:

```
(1, 224, 224, 3)
```

In [60]:

```
1 preds =model.predict(img_data)
```

```
1/1 [=====] - 0s 166ms/step
```

In [61]:

```
1 a=np.argmax(model.predict(img_data), axis=1)
```

1/1 [=====] - 0s 222ms/step

In [62]:

```
1 preds=np.argmax(preds, axis=1)
```

In [65]:

```
1 preds
```

Out[65]:

```
array([2], dtype=int64)
```

In [77]:

```
1 #COnverting the model to tensorflow lite for mobile app
```

In []:

```
1 #b. Mobil app model
```

In [73]:

```
1 import tensorflow as tf
2 from tensorflow.keras.models import load_model
3
4 # Load the best saved model from our last training
5 myModel = load_model('model_resnet50.h5')
```

In [74]:

```
1 # create a TFLiteConverter object from a TensorFlow Keras model
2 converter = tf.lite.TFLiteConverter.from_keras_model(myModel)
3
4 # converts a Keras model based on instance variable
5 myModel_tflite = converter.convert()
```

WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 5 of 53). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to: C:\Users\hbolla\AppData\Local\Temp\tmp4bq81o_w\assets

INFO:tensorflow:Assets written to: C:\Users\hbolla\AppData\Local\Temp\tmp4bq81o_w\assets

In [75]:

```
1 from pathlib import Path
2
3 # Save the model
4 tflite_model_file = Path('car_model_classifier.tflite')
5 tflite_model_file.write_bytes(myModel_tflite)
6
7 # with tf.io.gfile.GFile('clothing_classifier.tflite', mode='wb') as file:
8 #     file.write(myModel_tflite)
```

Out[75]:

96780488

```
1 # By Harsha Teja
```

In []:

```
1
```